INDIAN DATA CLUB

CODE BASICS

databricks

14 DAYS

AI CHALLENGE

DAY 13

Topic:

Model Comparison & Feature Engineering

Challenge:

1. Train 3 different models
2. Compare metrics in MLflow
3. Build Spark ML pipeline
4. Select best model

#DatabricksWithIDC

11:25 PM (6s) 3

```
#Loading Data
# Prepare data
df = spark.table("ecommerce.gold.product_metrics").toPandas()
if df.shape[0] == 0:
    print("No data available in ecommerce.silver.product_metrics.
    Cannot proceed with train/test split.")
else:
    X = df[["total_events"]]
    y = df["purchases"]
    X_train, X_test, y_train, y_test = train_test_split(X, y,
    test_size=0.2, random_state=42)
```

See performance (1)

df: pandas.core.frame.DataFrame = [product_id: object, brand: object ... 4 more fields]

X: pandas.core.frame.DataFrame = [total_events: int64]

X_test: pandas.core.frame.DataFrame = [total_events: int64]

X_train: pandas.core.frame.DataFrame = [total_events: int64]

11:34 PM (<1s) 4: Cell 4

```
models = {
    "LinearRegression": LinearRegression(),
    "DecisionTreeRegressor": DecisionTreeRegressor(max_depth = 5),
    "RandomForestRegressor": RandomForestRegressor
    (n_estimators=100, random_state=42)
}
```

Just now (3s) 7 Python

```python
from pyspark.ml import Pipeline
from pyspark.ml.feature import VectorAssembler
from pyspark.ml.regression import LinearRegression

assembler = vectorAssembler = VectorAssembler(inputCols=["total_events"], outputCol="features")
lr = LinearRegression(featuresCol="features", labelCol="purchases")
pipeline = Pipeline(stages=[assembler, lr])
spark_df =spark.table("ecommerce.gold.product_metrics").fillna(0)
train, test = spark_df.randomSplit([0.8, 0.2], seed=42)
pipeline_model = pipeline.fit(train)
# Evaluate the model
predictions = pipeline_model.transform(test)
display(predictions.select("product_id", "brand", "purchases", "prediction"))
```

See performance (1)

predictions: pyspark.sql.connect.dataframe.DataFrame
spark_df: pyspark.sql.connect.dataframe.DataFrame = [product_id: string, brand: string ... 4 more fields]
test: pyspark.sql.connect.dataframe.DataFrame = [product_id: string, brand: string ... 4 more fields]
train: pyspark.sql.connect.dataframe.DataFrame = [product_id: string, brand: string ... 4 more fields]

Table +

	product_id	brand	purchases	prediction
1	100000010	eksmo	0	-19.7436973545211...
2	100000014	eksmo	0	-20.3976444950013...
3	100000017	null	18	10.555853487728463

MLflow 3 is available!

Featuring unified ML and GenAI experiment tracking, improved model logging, prompt versior
beyond. Learn more

metrics.rmse < 1 and params.model = "tree"

Sort: Created

Group by

Run Name
RandomForestRegressor
DecisionTreeRegressor
LinearRegression
RandomForestRegressor
DecisionTreeRegressor
LinearRegression
LinearRegression
LinearRegression

Search metric charts

Model metrics (1)